

LRMoE RealData Preliminary

Spark Tseung

March 3, 2020

Introduction

This document is Part I of a demo series of the LRMoE (Logit-weighted Reduced Mixture-of-Experts) package on a real dataset. By analysing a French motor third-party liability insurance dataset in `CASdatasets`, we will demonstrate the fitting procedure, diagnostics, visualization and predictive functions of the LRMoE package. In this document, we focus on data preprocessing, exploratory analysis and parameter initialization.

Data Preprocessing

In this section, we first perform some exploratory analysis on the dataset. Description of variables are given in the `CASdatasets` user manual, and will not be repeated here. `freMTPLfreq` contains records of policy ID, policyholder information and number of claims, while `freMTPLsev` contains only policy ID and claim amounts for those who have claims. The `PolicyID` variable serves as a unique identifier that links these two datasets. We note that some policyholders have multiple claims, hence `freMTPLsev$PolicyID` contains duplicated records. We aggregate claim amounts by policy and merge these two datasets.

```
# Load data
data(freMTPLfreq, freMTPLsev)

# Aggregate claim amounts
sev.aggre = aggregate(freMTPLsev$ClaimAmount,
                      by = list(PolicyID = freMTPLsev$PolicyID), FUN = "sum")
colnames(sev.aggre)[2] = "ClaimAmount"

# Match two datasets by PolicyID. NA values correspond to no claims.
df.all = merge(freMTPLfreq, sev.aggre, by = "PolicyID", all = TRUE)
df.all$ClaimAmount[is.na(df.all$ClaimAmount)] = 0
df.all.pos = df.all[which(df.all$ClaimAmount>0),]

df.temp = data.matrix(df.all)
```

The response variables are `ClaimNb` and `ClaimAmount`. For simplicity, we exclude `Exposure` and `Density` from the covariates.

```
# Drop covariates: Exposure and Density
df.all = df.all[, -which(names(df.all) %in% c("Exposure", "Density"))]
df.all.pos = df.all.pos[, -which(names(df.all.pos) %in% c("Exposure", "Density"))]
```

Exploratory Analysis

We first calculate some summary statistics of `ClaimNb` and `ClaimAmount`. Less than 4% of policyholders have at least one claim, and the claim amount is right-skewed with a heavy tail.

```
sample.size = nrow(df.all)

# Proportion of zero claim
sum(df.all$ClaimNb==0)/sample.size

## [1] 0.9627513

sum(df.all$ClaimAmount==0)/sample.size

## [1] 0.9627513

# Other statistics
summary(df.all$ClaimAmount)

##      Min.   1st Qu.   Median     Mean   3rd Qu.     Max.
##      0.0      0.0      0.0     83.4      0.0 2036833.0

sd(df.all$ClaimNb[which(df.all$ClaimNb>0)])

## [1] 0.2314458

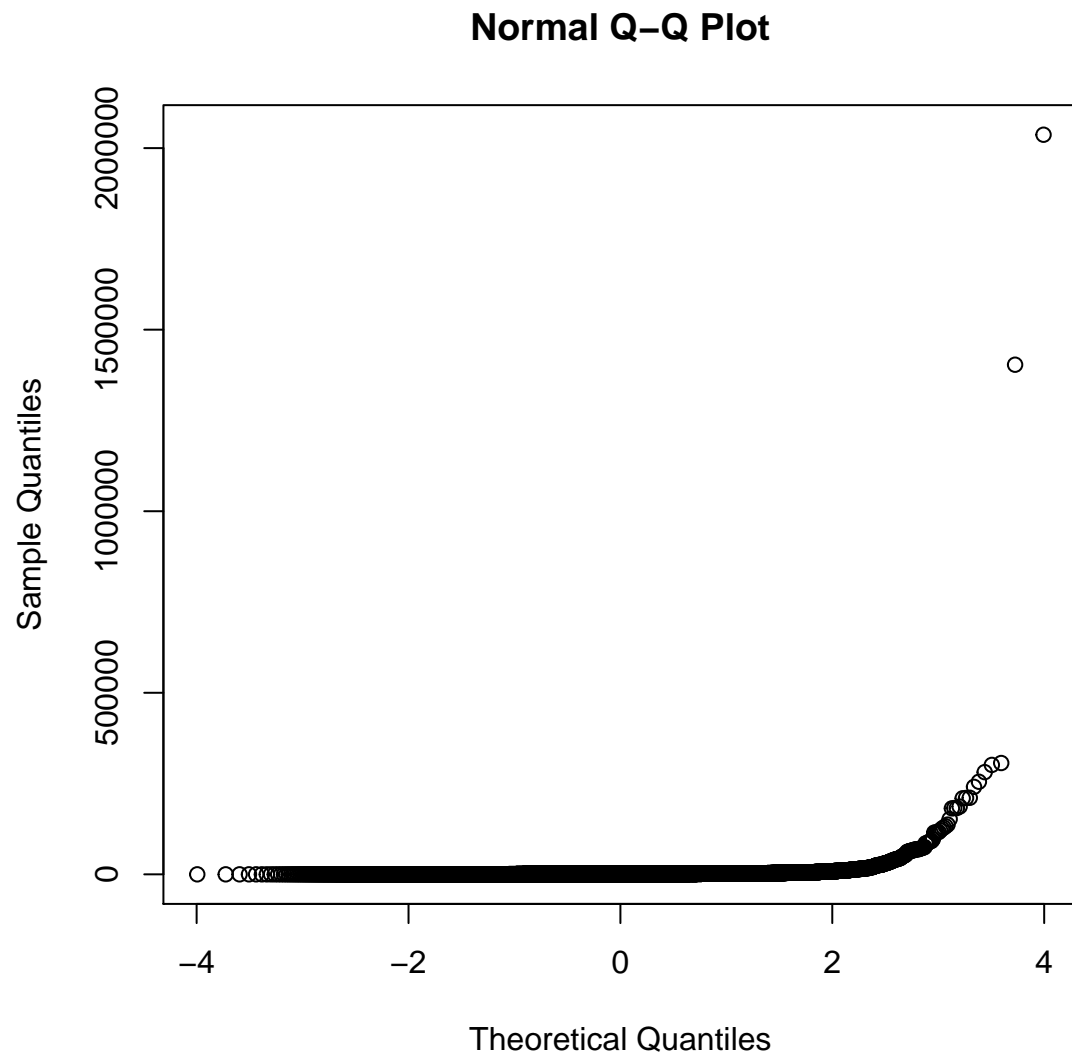
sd(df.all$ClaimAmount[which(df.all$ClaimAmount>0)])

## [1] 21612.28

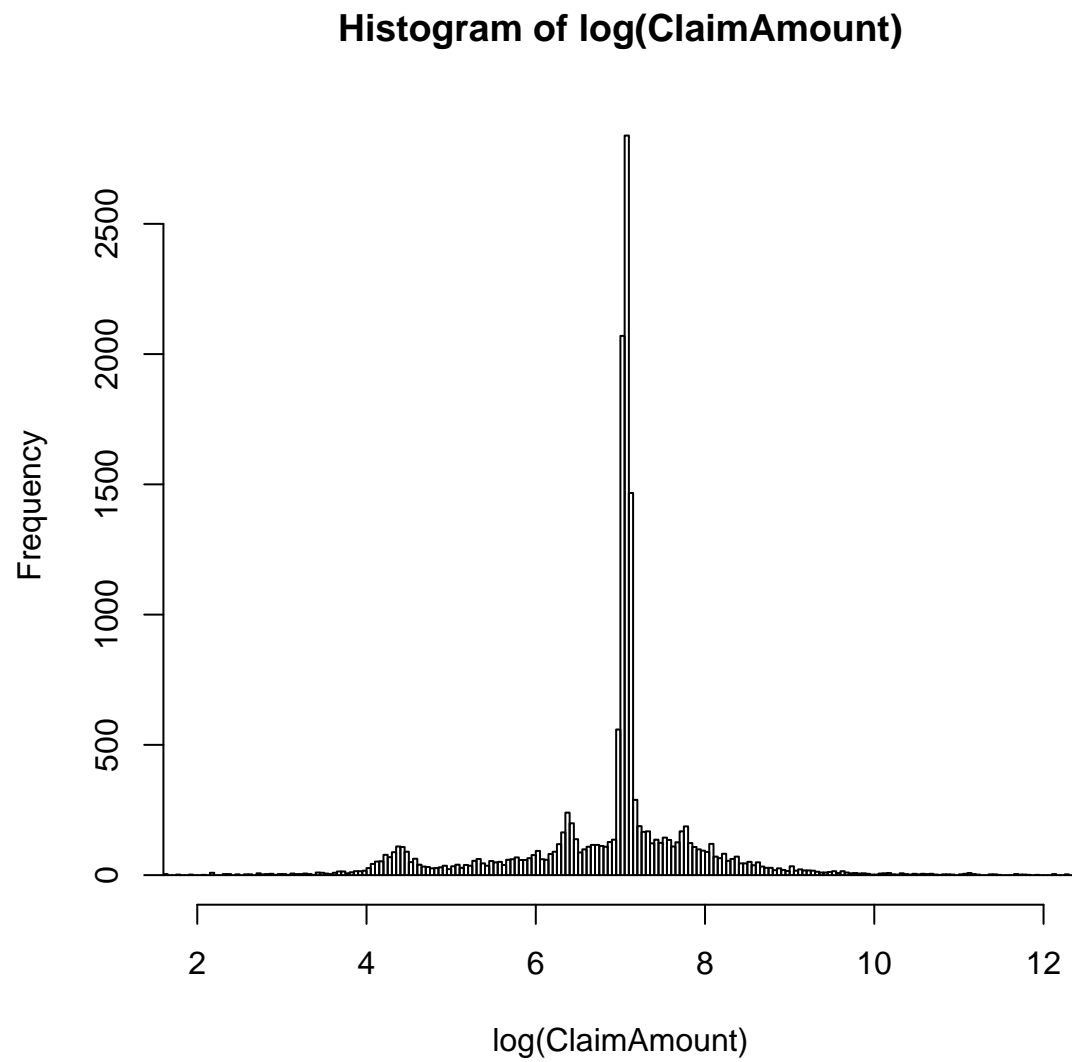
kurtosis(df.all$ClaimAmount[which(df.all$ClaimAmount>0)])

## [1] 6264.353
```

```
# QQ Plot  
qqnorm(df.all$ClaimAmount[which(df.all$ClaimAmount>0)])
```




```
# Histogram
hist(log(df.all$ClaimAmount[which(df.all$ClaimAmount>0 & df.all$ClaimAmount<500000)]),
     breaks = 200, probability = FALSE,
     xlim = c(2, 12),
     main = "Histogram of log(ClaimAmount)", xlab = "log(ClaimAmount)")
```

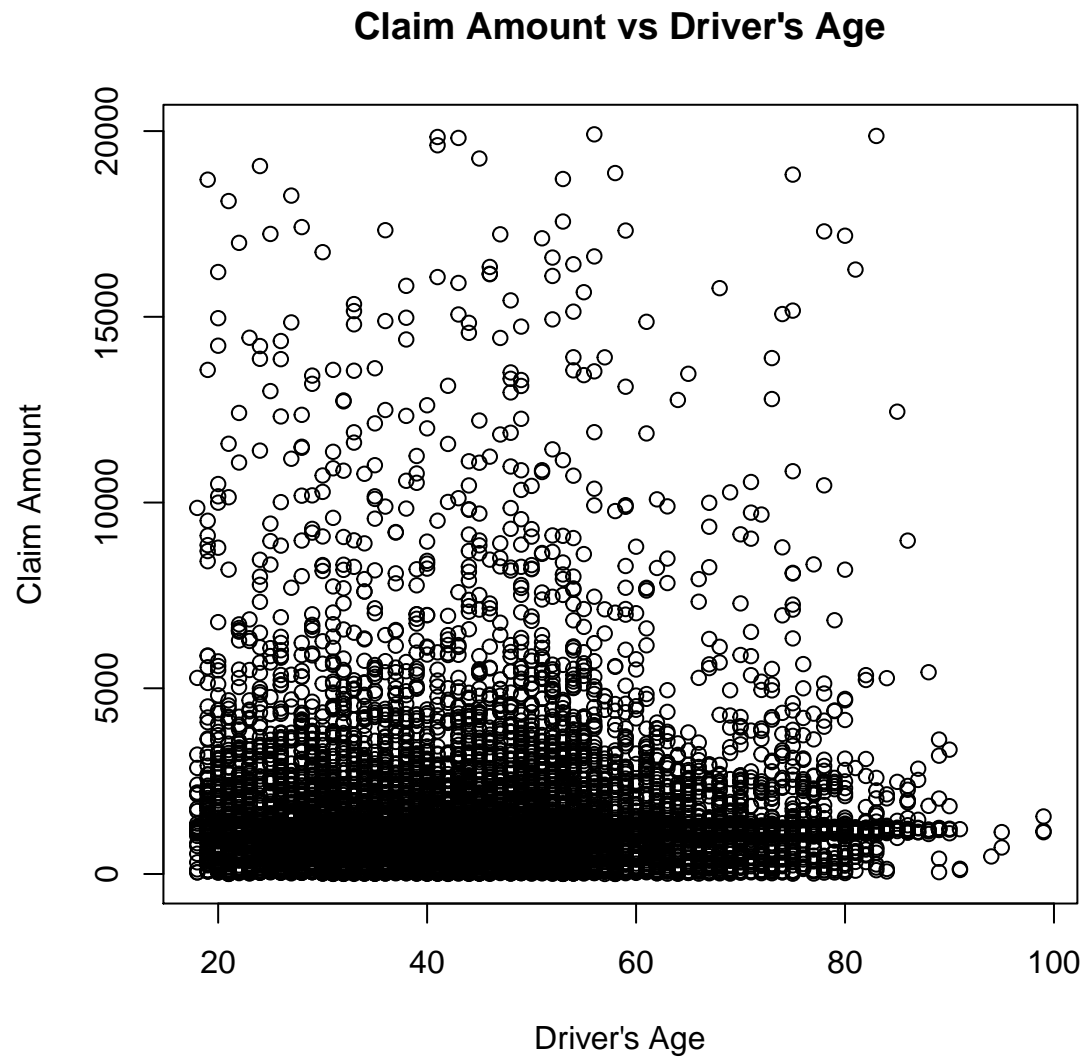


We can also plot the response variable against each of the covariates. For better presentation, only positive claim amounts less than 20,000 (or 5,000) are plotted.

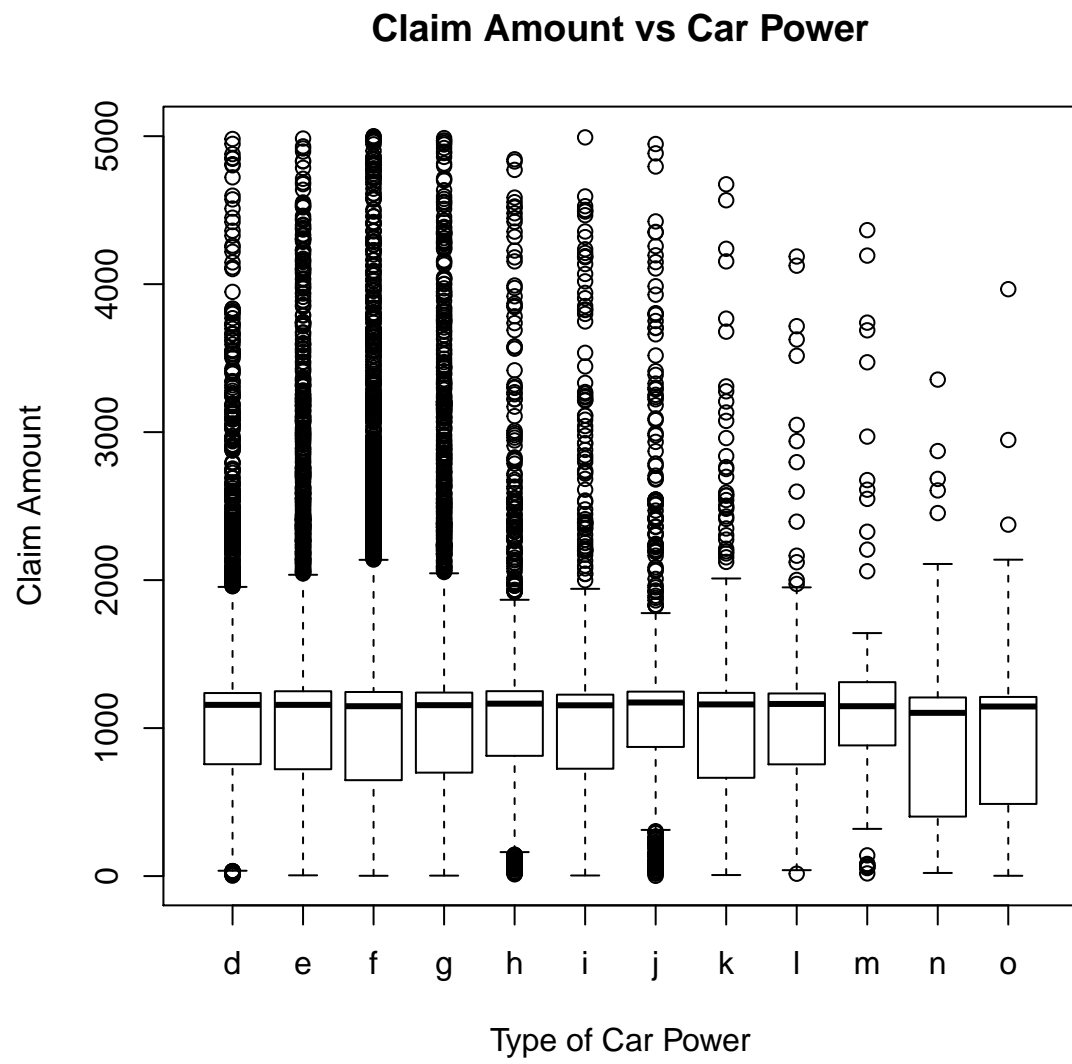
- Car's Age:



- Driver's Age:

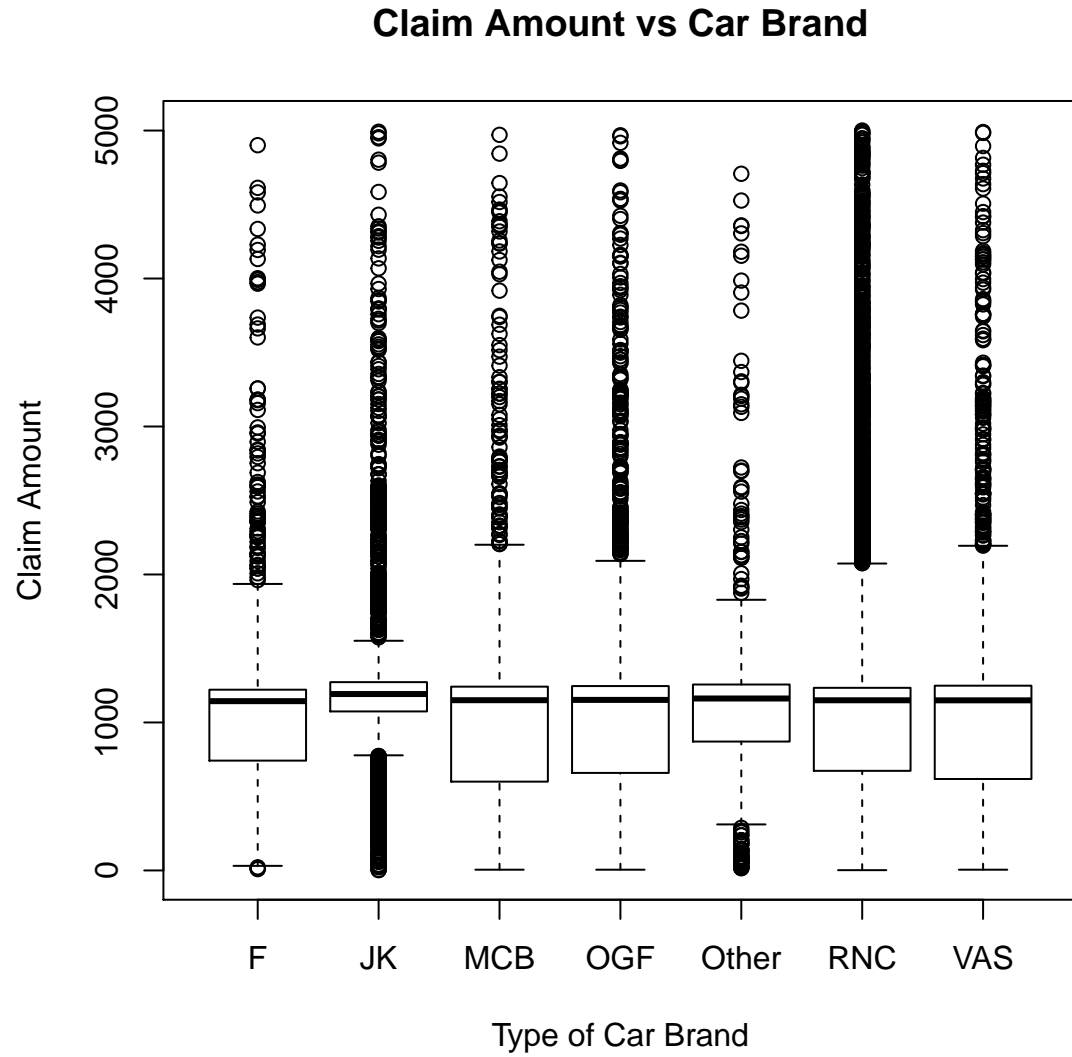


- Car Power:

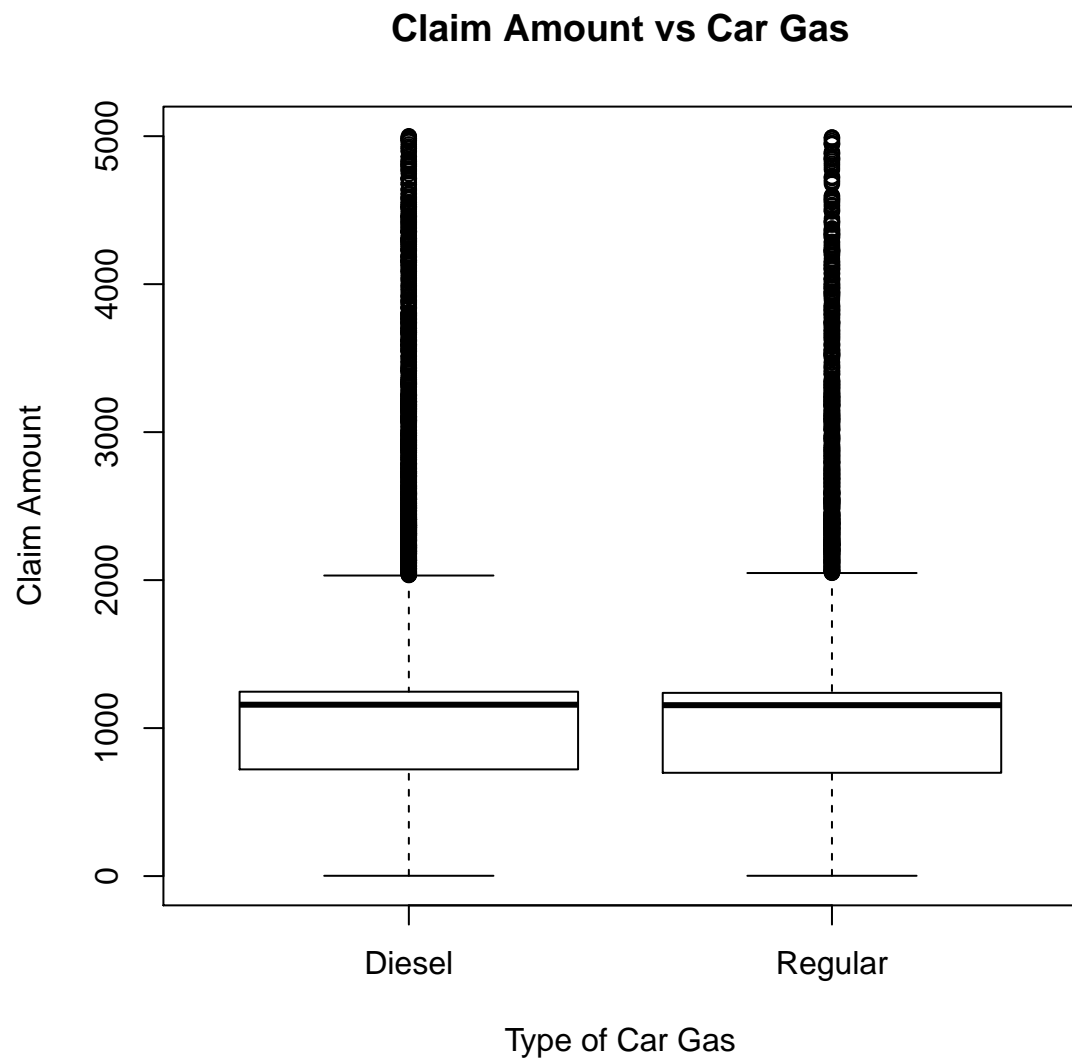


- Brand:

## [1] "Fiat"	"Japanese (except Nissan) or Korean"
## [3] "Mercedes, Chrysler or BMW"	"Opel, General Motors or Ford"
## [5] "other"	"Renault, Nissan or Citroen"
## [7] "Volkswagen, Audi, Skoda or Seat"	



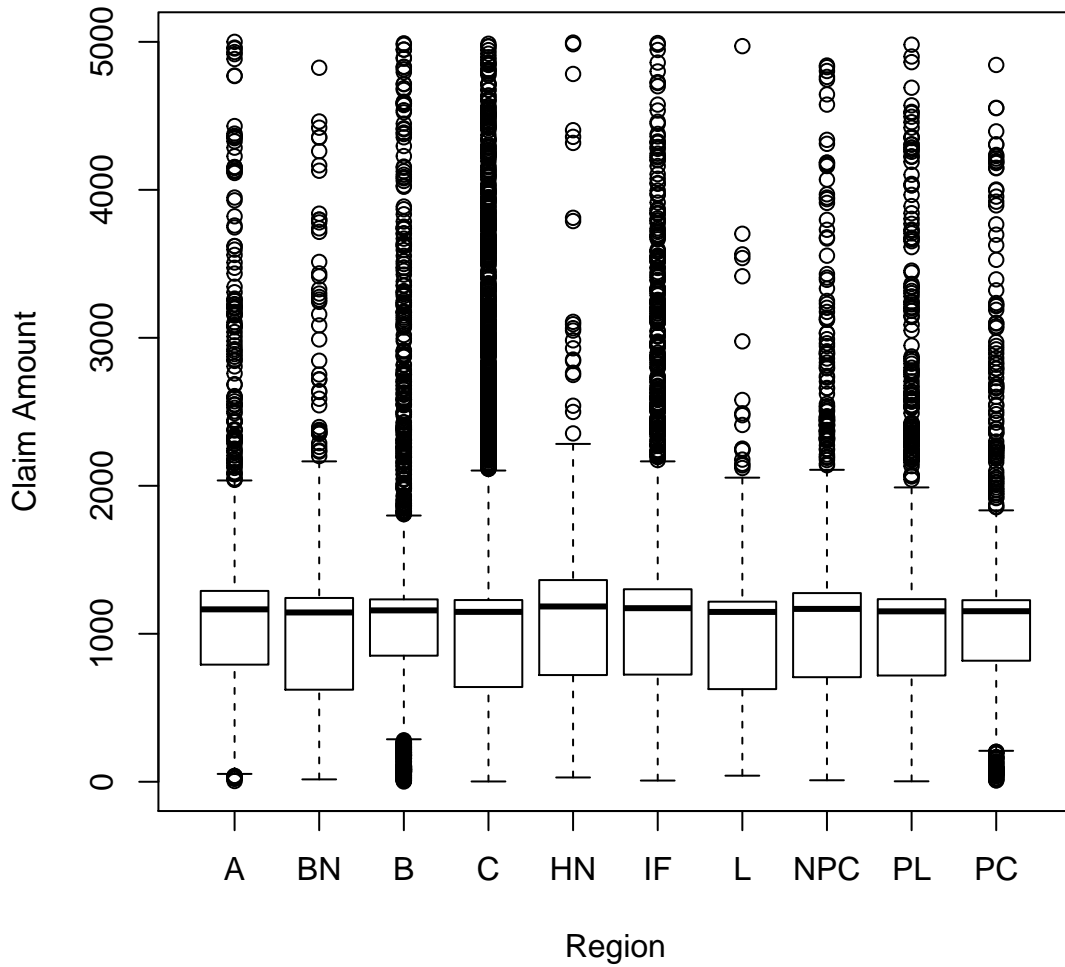
- Gas:



- Region:

```
## [1] "Aquitaine"          "Basse-Normandie"    "Bretagne"
## [4] "Centre"             "Haute-Normandie"    "Ile-de-France"
## [7] "Limousin"           "Nord-Pas-de-Calais" "Pays-de-la-Loire"
## [10] "Poitou-Charentes"
```

Claim Amount vs Region



Formatting Data

For illustration purposes, we will fit an LRMoE with one response variable `ClaimAmount` only. The following code shows how to format input matrices `Y` and `X` required by the LRMoE package.

```
# Make Y matrix
sample.size = nrow(df.all)
Y = matrix( c(rep(0, sample.size),      # = tl.1
              df.all$ClaimAmount,      # = yl.1
              df.all$ClaimAmount,      # = tu.1
              rep(Inf, sample.size)     # = tu.1
            ),
            ncol = 4, byrow = FALSE)

# Make X matrix
X.continuous = cbind(df.all$CarAge, df.all$DriverAge)
X.power = model.matrix(~df.all$Power, data = df.all)      # Default Power is 'd'
X.brand = model.matrix(~df.all$Brand, data = df.all)      # Default Brand is 'Fiat'
X.gas = model.matrix(~df.all$Gas, data = df.all)          # Default Gas is 'Diesel'
X.region = model.matrix(~df.all$Region, data = df.all)    # Default Region is 'Aquitaine'
X = matrix(cbind(rep(1, sample.size), # Intercept
                 X.continuous, X.power[,-1], X.brand[,-1], X.gas[,-1], X.region[,-1]),
           nrow = sample.size, byrow = FALSE)

colnames(X) = c("Intercept", "CarAge", "DriverAge",
               "Powere", "Powerf", "Powerg", "Powerh", "Poweri", "Powerj",
               "Powerk", "Powerl", "Powerm", "Powern", "Powero",
               "BrandJK", "BrandMCB", "BrandOGF", "BrandOther", "BrandRNC", "BrandVAS",
               "GasRegular",
               "RegionBN", "RegionB", "RegionC", "RegionHN", "RegionIF",
               "RegionL", "RegionNPC", "RegionPL", "RegionPC")
```

We will save the data for the fitting procedures later.

```
save(X, file = "X.Rda")
save(Y, file = "Y.Rda")
```

Parameter Initialization

The LRMoE fitting function also requires initialization of `n.comp` (number of latent components), `comp.dist` (component distributions by dimension and by component), `zero.init` (zero inflation) and `params.init` (initialization of component distribution parameters).

Since component distributions included in LRMoE are all uni-modal, a good starting point is to observe the numbers of components is to count the number of peaks in the previous histogram of data. We will consider 3~6 latent components, each with 5 combinations of component distributions. For each case, we use k-means clustering and matching of moments to roughly choose initial parameters.

```
# Drop response: ClaimNb
df.all = df.all[, -which(names(df.all) %in% c("ClaimNb"))]
df.all.pos = df.all.pos[, -which(names(df.all.pos) %in% c("ClaimNb"))]

# Normalize data
df.all.norm = df.all
df.all.norm = df.all.norm[, -which(names(df.all.norm) %in% c("PolicyID", "ClaimNb"))]

df.all.norm$CarAge =
```

```
(df.all.norm$CarAge - mean(df.all.norm$CarAge))/sd(df.all.norm$CarAge)
df.all.norm$DriverAge =
  (df.all.norm$DriverAge - mean(df.all.norm$DriverAge))/sd(df.all.norm$DriverAge)
df.all.norm$ClaimAmount =
  (df.all.norm$ClaimAmount - mean(df.all.norm$ClaimAmount))/sd(df.all.norm$ClaimAmount)
```

```
head(df.all.norm)
```

```
##      Power      CarAge      DriverAge      Brand      Gas
## 1      g -1.3070259  0.04746775 Japanese (except Nissan) or Korean Diesel
## 2      g -1.3070259  0.04746775 Japanese (except Nissan) or Korean Diesel
## 3      f -0.9599851 -0.51087486 Japanese (except Nissan) or Korean Regular
## 4      f -0.9599851 -0.51087486 Japanese (except Nissan) or Korean Regular
## 5      g -1.3070259 -0.30149638 Japanese (except Nissan) or Korean Diesel
## 6      g -1.3070259 -0.30149638 Japanese (except Nissan) or Korean Diesel
##
##      Region ClaimAmount
## 1      Aquitaine -0.01989646
## 2      Aquitaine -0.01989646
## 3 Nord-Pas-de-Calais -0.01989646
## 4 Nord-Pas-de-Calais -0.01989646
## 5 Pays-de-la-Loire -0.01989646
## 6 Pays-de-la-Loire -0.01989646
```

```
summary(df.all.norm)
```

```
##      Power      CarAge      DriverAge
## f      :95718  Min.    :-1.30703  Min.    :-1.90673
## g      :91198  1st Qu.: -0.78646  1st Qu.: -0.79005
## e      :77022  Median  :-0.09238  Median  :-0.09212
## d      :68014  Mean    : 0.00000  Mean    : 0.00000
## h      :26698  3rd Qu.: 0.77522  3rd Qu.: 0.60581
## j      :18038  Max.    :16.04501  Max.    : 3.74649
## (Other):36481
##
##      Brand      Gas
## Fiat      : 16723  Diesel :205945
## Japanese (except Nissan) or Korean: 79060  Regular:207224
## Mercedes, Chrysler or BMW      : 19280
## Opel, General Motors or Ford    : 37402
## other      : 9866
## Renault, Nissan or Citroen      :218200
## Volkswagen, Audi, Skoda or Seat : 32638
##
##      Region      ClaimAmount
## Centre      :160601  Min.    : -0.0199
## Ile-de-France : 69791  1st Qu.: -0.0199
## Bretagne     : 42122  Median  : -0.0199
## Pays-de-la-Loire : 38751  Mean    : 0.0000
## Aquitaine     : 31329  3rd Qu.: -0.0199
## Nord-Pas-de-Calais: 27285  Max.    :485.8049
## (Other)      : 43290
```

The LRMoe package contains a Clustered Method of Moments initialization function which is used in conjunction of kmeans. We look at the 3-component case in detail and skip the rest.

3 Latent Components

```
set.seed(7777) # For reproducible results

dim.m = 1
n.comp = 3

norm.km.cluster.3 = kmeans(data.matrix(df.all.norm), n.comp)
norm.init.analysis.3 = cluster.mm.severity(df.all$ClaimAmount, norm.km.cluster.3$cluster)
```

The returned list `norm.init.analysis.3` contains cluster proportion (of the entire population), zero inflation, summary statistics and parameter initializations for all selection of component distributions. The user can then choose which distributions to use. As a general rule of thumb, initialization with very extreme parameter values should be avoided.

For example, the initialization of the first component is as follows.

```
norm.init.analysis.3[[1]]

## $cluster.prop
## [1] 0.207397
##
## $zero.prop
## [1] 0.9625394
##
## $mean.pos
## [1] 1736.804
##
## $var.pos
## [1] 14778491
##
## $cv.pos
## [1] 2.213422
##
## $skew.pos
## [1] 10.50369
##
## $kurt.pos
## [1] 143.7694
##
## $gamma.init
##      shape      scale
## 0.2041134 8509.0148004
##
## $lnorm.init
## meanlog  sdlog
## 6.572390 1.332225
##
## $invgauss.init
##      mean      scale
## 1736.804 354.505
##
## $weibull.init
##      shape      scale
## 3.000 2605.206
```

```
##
## $burr.init
##   shape1   shape2   scale
##     2.00     5.00 12770.62
```

4~6 Latent Components

```
set.seed(7777) # For reproducible results

dim.m = 1
n.comp = 4

norm.km.cluster.4 = kmeans(data.matrix(df.all.norm), n.comp)
norm.init.analysis.4 = cluster.mm.severity(df.all$ClaimAmount, norm.km.cluster.4$cluster)

set.seed(7777) # For reproducible results

dim.m = 1
n.comp = 5

norm.km.cluster.5 = kmeans(data.matrix(df.all.norm), n.comp)
norm.init.analysis.5 = cluster.mm.severity(df.all$ClaimAmount, norm.km.cluster.5$cluster)

set.seed(7777) # For reproducible results

dim.m = 1
n.comp = 6

norm.km.cluster.6 = kmeans(data.matrix(df.all.norm), n.comp)
norm.init.analysis.6 = cluster.mm.severity(df.all$ClaimAmount, norm.km.cluster.6$cluster)
```

We will save all initalizations for use in the next step.

```
save(norm.init.analysis.3, file = "RealDataInit3.Rda")
save(norm.init.analysis.4, file = "RealDataInit4.Rda")
save(norm.init.analysis.5, file = "RealDataInit5.Rda")
save(norm.init.analysis.6, file = "RealDataInit6.Rda")
```